



SCHOOL OF COMPUTING AND INFORMATION TECHNOLOGY

B22CI0501 – MACHINE LEARNING

PREDICTION OF DIABETES PAITENTS
USING DIABETES DATASET

SUBMITTED BY

NAME: TARUN KUMAR S

SRN: R22EA063

SEMESTER & BRANCH: 5th SEM AIML B

DATE OF SUBMISSION:

SIGNATURE:

VERIFIED BY:

MARKS:

SIGNATURE OF FACULTY:

QUESTIONS TO BE ANSWERED:

1. What is the nature of the dataset (structured, unstructured, or semi-structured)?
 - The dataset is **structured**, with clear columns representing variables (features) and the "Survived" as the target label. It is a **labelled** dataset.
2. What is the source of the dataset (public, proprietary, or in-house)?
 - The Titanic dataset primarily comes from passenger records of the RMS Titanic, a British passenger liner that sank on its maiden voyage in 1912.
3. What is the size of the dataset in terms of samples or instances before applying PCA?
 - The Titanic dataset from Kaggle typically has:
 - 891 samples (instances) in the training set.
 - 418 samples in the test set.
4. Are the labels accurate and meaningful for the intended task?
 - The labels in the "Survived" column are binary (0 for Not Survived, 1 for Survived). They appear meaningful for the task of predicting.
5. How consistent is the labeling across the dataset?
 - Since the dataset is structured and fully populated, the labeling seems consistent across all instances.

6. Are there missing values or corrupted data points in the dataset?

- No missing values are detected in this dataset. However, some values (like zero values for Age and Cabin) may need to be addressed as potential data issues or outliers.

7. Are there any legal or ethical issues related to using this dataset?

- No, there is no legal or ethical issues related to using this dataset.

8. What is the class distribution, and is it balanced or imbalanced?

- The target class distribution (Survived) will need to be checked to see if it's balanced or imbalanced.

- Based on analysis, The dataset is slightly imbalanced:

Class 0 (Not Survived): 60%

Class 1 (Survived): 40%

CODE OF THE ML MODEL:

#Step 1: Data Pre-processing step

#importing libraries

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score, confusion_matrix,
ConfusionMatrixDisplay, classification_report

from sklearn.preprocessing import StandardScaler

#importing dataset

#Load the data

train_data = pd.read_csv('../dataset/raw/train.csv')

test_data = pd.read_csv('../dataset/raw/test.csv')

test_data_survived = pd.read_csv('../dataset/raw/gender_submission.csv')

```
# Handle missing values
```

```
train_data.fillna(method='ffill', inplace=True)
```

```
test_data.fillna(method='ffill', inplace=True)
```

```
# Convert categorical variables to numerical
```

```
train_data = pd.get_dummies(train_data, columns=['Sex', 'Embarked'],  
drop_first=True)
```

```
test_data = pd.get_dummies(test_data, columns=['Sex', 'Embarked'],  
drop_first=True)
```

```
#Check the data
```

```
print("Train Data:")
```

```
print(train_data.head())
```

```
print("\nTest Data:")
```

```
print(test_data.head())
```

```
print(test_data_survived["Survived"].head())
```

Tarin Data:

	PassengerId	Survived	Pclass \
0	1	0	3
1	2	1	1
2	3	1	3
3	4	1	1
4	5	0	3

	Name	Age	SibSp	Parch \
0	Braund, Mr. Owen Harris	22.0	1	0
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	38.0	1	0
2	Heikkinen, Miss. Laina	26.0	0	0
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	35.0	1	0
4	Allen, Mr. William Henry	35.0	0	0

	Ticket	Fare	Cabin	Sex_male	Embarked_Q	Embarked_S
0	A/5 21171	7.2500	NaN	True	False	True
1	PC 17599	71.2833	C85	False	False	False
2	STON/O2. 3101282	7.9250	C85	False	False	True
3	113803	53.1000	C123	False	False	True
4	373450	8.0500	C123	True	False	True

Test Data:

	PassengerId	Pclass	Name	Age \
0	892	3	Kelly, Mr. James	34.5
1	893	3	Wilkes, Mrs. James (Ellen Needs)	47.0
2	894	2	Myles, Mr. Thomas Francis	62.0
3	895	3	Wirz, Mr. Albert	27.0
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	22.0

	SibSp	Parch	Ticket	Fare	Cabin	Sex_male	Embarked_Q	Embarked_S
0	0	0	330911	7.8292	NaN	True	True	False
1	1	0	363272	7.0000	NaN	False	False	True
2	0	0	240276	9.6875	NaN	True	True	False
3	0	0	315154	8.6625	NaN	True	False	True
4	1	1	3101298	12.2875	NaN	False	False	True
0	0							
1	1							
2	0							
3	0							
4	1							

```
# Select features and target

X_train = train_data.drop(['Survived', 'Name', 'Ticket', 'Cabin'], axis=1)

y_train = train_data['Survived']

X_test = test_data.drop(['Name', 'Ticket', 'Cabin'], axis=1)


#Fit the train and test data

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test).


#Displaying the data

print("Training set of Independent after features scaling:")

print(X_train[:5])

print("\nTraining set of Dependent:")

print(y_train[:5])

print("\nTesting set of Independent after features scaling:")

print(X_test[:5])

print("\ntesting set of Dependent:")

print(y_test[:5])
```

```

Training set of Independent variables after features scaling:
[[-1.73010796  0.82737724 -0.52119766  0.43279337 -0.47367361 -0.50244517
  0.73769513 -0.30974338  0.61930636]
 [-1.72622007 -1.56610693  0.57872934  0.43279337 -0.47367361  0.78684529
 -1.35557354 -0.30974338 -1.61470971]
 [-1.72233219  0.82737724 -0.24621591 -0.4745452  -0.47367361 -0.48885426
 -1.35557354 -0.30974338  0.61930636]
 [-1.71844431 -1.56610693  0.37249302  0.43279337 -0.47367361  0.42073024
 -1.35557354 -0.30974338  0.61930636]
 [-1.71455642  0.82737724  0.37249302 -0.4745452  -0.47367361 -0.48633742
  0.73769513 -0.30974338  0.61930636]]

```

```

Training set of Dependent variables:

```

```

0  0
1  1
2  1
3  1
4  0

```

```

Name: Survived, dtype: int64

```

```

Testing set of Independent variables after features scaling:

```

```

[[ 1.73399584  0.82737724  0.33812031 -0.4745452  -0.47367361 -0.49078316
  0.73769513  3.22847904 -1.61470971]
 [ 1.73788372  0.82737724  1.19743827  0.43279337 -0.47367361 -0.50747884
 -1.35557354 -0.30974338  0.61930636]
 [ 1.74177161 -0.36936484  2.22861983 -0.4745452  -0.47367361 -0.45336687
  0.73769513  3.22847904 -1.61470971]
 [ 1.74565949  0.82737724 -0.17747047 -0.4745452  -0.47367361 -0.47400493
  0.73769513 -0.30974338  0.61930636]
 [ 1.74954737  0.82737724 -0.52119766  0.43279337  0.76762988 -0.40101668
 -1.35557354 -0.30974338  0.61930636]]

```

```

Testing set of Dependent variables:

```

```

0  0
1  1
2  0
3  0
4  1


```


#Step 2: Training the model

#Fitting Logistic Regression to the training set

```
model = LogisticRegression(max_iter=10, random_state=7)
```

```
model.fit(X_train, y_train)
```



#Step 3: Predicting the test set result

Predict on the test data

```
test_predictions = model.predict(X_test)
```

```
print(test_predictions)
```

```
[0000101010001011001100111010000000101
1000001100011001100000100011110111011
1101010000001110101010101010001000000
1111001111010010100000000001001000000
1010010011011010010011000001101100101
0100000000011011001011010000100101010
1010010001000000111100001011101000001
0001100001000110100001011100000110000
1000000011000000011100000000101000110
1000000000101010110001010010110100110
0100111000001101000011000101001011000
01111101000]
Test Accuracy: 0.937799043062201
```

#Step 4: Creating the confusion matrix and checking other performance metrics

Calculate the accuracy

```
accuracy = accuracy_score(y_test, test_predictions)
```

```
print(f'Test Accuracy: {accuracy}')
```

Calculate the confusion matrix

```
conf_matrix = confusion_matrix(y_test, test_predictions)
```

```
print("\nConfusion Matrix:")
```

```
print(conf_matrix)
```

Display the classification report

```
class_report = classification_report(y_test, test_predictions, target_names=['Not  
Survived', 'Survived'], output_dict=True)
```

```
class_report_df = pd.DataFrame(class_report).transpose()
```

```
class_report_df = class_report_df.round(2)
```

```
print("\nFormatted Classification Report:")
```

```
print(class_report_df)
```

```
# Display the confusion matrix using matplotlib
```

```
disp = ConfusionMatrixDisplay(confusion_matrix=conf_matrix)
```

```
disp.plot(cmap=plt.cm.Blues)
```

```
plt.show()
```


Test Accuracy: 0.937799043062201

Confusion Matrix:

```
[[250 16]
 [ 10 142]]
```

Formatted Classification Report:

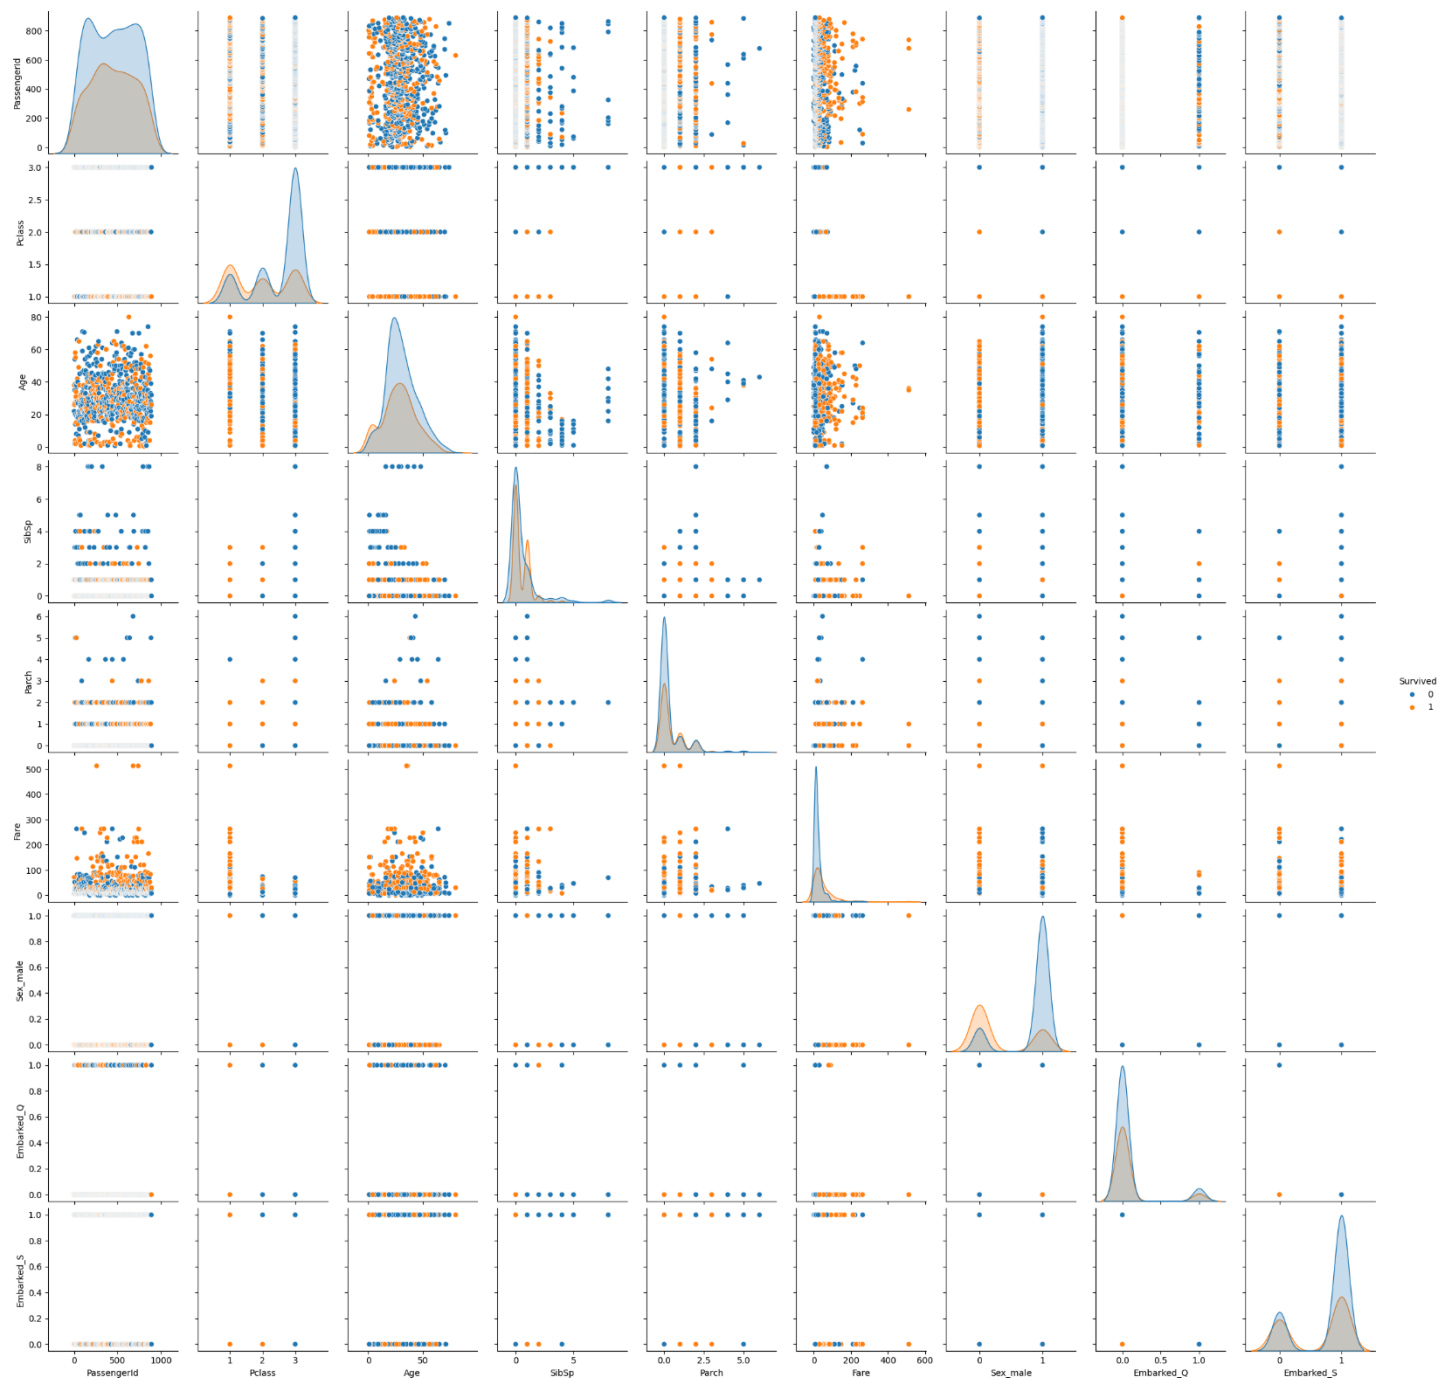
	precision	recall	f1-score	support
Not Survived	0.96	0.94	0.95	266.00
Survived	0.90	0.93	0.92	152.00
accuracy	0.94	0.94	0.94	0.94
macro avg	0.93	0.94	0.93	418.00
weighted avg	0.94	0.94	0.94	418.00



Pair plot to visualize pairwise relationships between features

```
sns.pairplot(train_data, hue='Survived', diag_kind='kde')
```

```
plt.show()
```



Predicting the output based on user input

```
#Predicting the output based on user input
```

```
import joblib
```

```
import numpy as np
```

```
import os
```

```
from sklearn.preprocessing import StandardScaler
```

```
# Function to load all models from a specified folder
```

```
def load_models_from_folder(folder_path):
```

```
    models = { }
```

```
    for filename in os.listdir(folder_path):
```

```
        if filename.endswith('.pkl'):
```

```
            model_name = filename.split('.')[0]
```

```
            model_path = os.path.join(folder_path, filename)
```

```
            models[model_name] = joblib.load(model_path)
```

```
    return models
```

```
# Function to take input for the Titanic dataset
```

```
def get_titanic_input():
```

```
    Pclass = int(input("Enter Pclass (1, 2, or 3): "))
```

```
    Sex = input("Enter Sex (male or female): ")
```

```
    Age = float(input("Enter Age: "))
```

```
    SibSp = int(input("Enter number of siblings/spouses aboard: "))
```

```
    Parch = int(input("Enter number of parents/children aboard: "))
```

```
    Fare = float(input("Enter Fare: "))
```

```
    Embarked = input("Enter Embarked (C, Q, or S): ")
```

```
# Convert categorical variables to numerical
Sex = 1 if Sex == 'male' else 0
Embarked_C = 1 if Embarked == 'C' else 0
Embarked_Q = 1 if Embarked == 'Q' else 0
Embarked_S = 1 if Embarked == 'S' else 0

# Create input array
input_data = np.array([[Pclass, Sex, Age, SibSp, Parch, Fare, Embarked_C,
Embarked_Q, Embarked_S]])

print("Pclass:", Pclass)
print("Sex:", Sex)
print("Age:", Age)
print("Siblings/Spouses:", SibSp)
print("Parents/Children:", Parch)
print("Fare:", Fare)
print("Embarked_C:", Embarked_C)
print("Embarked_Q:", Embarked_Q)
print("Embarked_S:", Embarked_S)
print("\n")
return input_data

# Load models from the specified folder
folder_path = './models' # Replace with the path to your models folder
models = load_models_from_folder(folder_path)

# Get input data
input_data = get_titanic_input()
```

```
survive_dict = {  
    0: "Not Survived",  
    1: "Survived"  
}  
  
# Predict using each model and print the outputs  
for model_name, model in models.items():  
    output = model.predict(input_data)  
    print(f'Output from {model_name}: {survive_dict[output[0]]}')
```

```
Pclass: 3  
Sex: 1  
Age: 35.0  
Siblings/Spouses: 0  
Parents/Children: 0  
Fare: 8.05  
Embarked_C: 0  
Embarked_Q: 0  
Embarked_S: 1  
  
Output from decision_tree: Not Survived  
Output from gradient_boosting: Survived  
Output from k-nearest_neighbors: Not Survived  
Output from logistic_regression: Not Survived  
Output from logistic_regression_model_main: Not Survived  
Output from naive_bayes: Survived  
Output from neural_network_(mlp): Not Survived  
Output from random_forest: Not Survived  
Output from support_vector_machine: Not Survived
```

```
#Using same dataset on different ML Algorithms

import matplotlib.pyplot as plt
import seaborn as sns
import joblib

from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score


#Defining different models to a Dictionary
models = {
    "Support Vector Machine": SVC(kernel='linear', random_state=0),
    "K-Nearest Neighbors": KNeighborsClassifier(),
    "Decision Tree": DecisionTreeClassifier(random_state=0),
    "Logistic Regression": LogisticRegression(random_state=0),
    "Random Forest": RandomForestClassifier(random_state=0),
    "Gradient Boosting": GradientBoostingClassifier(random_state=0),
    "Neural Network (MLP)": MLPClassifier(random_state=0),
    "Naive Bayes": GaussianNB(),
}
```



```

#Storing the accuracies of every model
results = []
for model_name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"{model_name} Accuracy: {accuracy * 100:.2f}%")
    results.append([model_name, accuracy])
    file_name = model_name.lower().replace(" ", "_") + ".pkl"
    joblib.dump(model, f"../models/{file_name}")

# Create a DataFrame for the results
results_df = pd.DataFrame(results, columns=["Model", "Accuracy"])

#Plotting different models accuracy
plt.figure(figsize=(12, 6)) # Adjust figure size as needed
ax = sns.barplot(x="Model", y="Accuracy", data=results_df)
plt.title("Comparison of Model Accuracies")
plt.xlabel("Model")
plt.ylabel("Accuracy")
plt.ylim(0, 1) # Set y-axis limits to 0-1 for accuracy
plt.xticks(rotation=45, ha="right") # Rotate x-axis labels for better readability

# Add text labels on the bars
for p in ax.patches:
    ax.annotate(f'{p.get_height():.2f}',
                (p.get_x() + p.get_width() / 2., p.get_height()),

```

```
ha='center', va='center',  
xytext=(0, 9),  
textcoords='offset points')
```

```
plt.tight_layout()
```

```
plt.show()
```

```
Support Vector Machine Accuracy: 100.00%  
K-Nearest Neighbors Accuracy: 81.34%  
Decision Tree Accuracy: 77.03%  
Logistic Regression Accuracy: 93.78%  
Random Forest Accuracy: 80.38%  
Gradient Boosting Accuracy: 82.78%  
Neural Network (MLP) Accuracy: 76.32%  
Naive Bayes Accuracy: 91.39%
```

